

VISHMITHA POOJARY

Mangalore, Karnataka +91 9110833308 ✉ vishuvishmitha84@gmail.com LinkedIn GitHub

PROFESSIONAL SUMMARY

I build backend systems, full-stack applications, and AI products. To ship features faster, I use AI IDEs well beyond simple autocomplete. Before I start coding, I set up an AGENTS.md file to give the AI exact project context, which helps me map out a solid implementation plan. I also automate repetitive tasks by packaging them into SKILL.md files. By treating AI as a highly capable assistant rather than a crutch (planning first, setting strict context, and offloading the busywork), I significantly boost my output and deliver high-quality code at a much faster pace.

EDUCATION

Course	Institution	Year	CGPA / %
MCA	St Aloysius Institute of Management	2024–2026	9.00
BCA	SDM College of Business Management	2021–2024	9.15

SKILLS

AI Dev Workflow: Codex (AGENTS.md, SKILL.md, plan mode, verification gates), Claude, Gemini, ChatGPT — structured prompting, AI-assisted debugging, prompt engineering

AI & Agents: LangGraph, Gemini API, OpenAI API, Agentic Function Calling, Multi-Agent Orchestration, Pydantic

Backend: Python, FastAPI, Node.js, TypeScript, Express.js, Hono.js **Frontend:** React.js, Next.js, TailwindCSS

Databases & Tools: PostgreSQL, MongoDB, Drizzle ORM, SQLAlchemy, Zod, JWT, Docker, Git, WebSockets

EXPERIENCE

Software Engineer — Trikon Software Labs <i>Mangalore, India</i>	Jul 2026 – Present
– Build backend services for a live enterprise voice AI platform — Node.js, TypeScript, Express, PostgreSQL, WebSockets — integrating Google Gemini Live API with telephony providers (Plivo, Vobiz) for real-time AI phone calls – Replaced static pre-recorded audio with a live multilingual AI voice engine supporting 30+ languages; built audio filters that stop the AI getting confused by background noise, eliminating hallucinations and runaway calls – Rebuilt the outbound dialer from polling to event-driven architecture with smart traffic pacing so it handles high-volume campaigns without hitting telecom rate limits; executed a zero-data-loss live production database migration (Supabase → Neon PostgreSQL) – Use Codex daily with AGENTS.md for persistent project context, SKILL.md for reusable workflows, and plan mode before every significant feature — when AI produces repeated errors, I trace the actual execution flow and update AGENTS.md so the mistake doesn't happen again	
Software Developer Intern — TechBrew Solutions <i>Mangalore, India</i>	Jan 2026 – Jun 2026
– Hospital Laundry Management System: Built production REST APIs end-to-end from OpenAPI contracts — schema design, PostgreSQL migrations (Drizzle ORM), Zod validation, business logic, and React frontend integration for a live hospital client – AI Multi-Agent Trading System: Worked on a 4-agent LangGraph system (Supervisor, StockNews, TechnicalAnalysis, TradeBrain) — used plan mode and Claude to design the agent architecture upfront; engineered structured LLM prompts with event taxonomy and Pydantic-validated JSON output, iteratively comparing Gemini and GPT to find which handled edge cases better	

PROJECTS

AI Hospital Voice Receptionist (Deployed) FastAPI, PostgreSQL, Gemini, LiveKit, Deepgram, ElevenLabs	2026
– Built and deployed an AI voice agent for a healthcare facility that handles appointment booking, rescheduling, and cancellations in English and Hindi over a real phone call — integrated LiveKit (WebRTC), Deepgram (speech-to-text), and ElevenLabs (text-to-speech) with Gemini function calling against live PostgreSQL data – Used Codex with a project AGENTS.md to scaffold the agentic architecture; when LiveKit integration kept failing, traced the actual audio pipeline rather than re-prompting, identified the root cause, and fixed it properly	
AI Interview Platform Backend (Ongoing) FastAPI, PostgreSQL, SQLAlchemy, JWT, LLM APIs	2026
– Used plan mode to design the full API structure before writing any code — 9 secured endpoints covering auth, adaptive interview sessions, LLM-driven question generation, PDF resume parsing, weak-skill tracking, and leaderboard analytics	
AI Multi-Agent Intraday Trading System Python, LangGraph, Gemini API, OpenAI API, Pydantic	2026
– Engineered the LLM news ranker with structured prompt engineering — event taxonomy (earnings, M&A, regulatory, debt restructuring), confidence/impact/freshness scoring, Pydantic-validated JSON output; found and fixed a cross-batch deduplication bug by tracing the full pipeline rather than guessing; built retry logic with exponential backoff and per-call token tracking	